



# DATABASE SYSTEMS

CS – 355/CE – 373

Instructor: Maria N. Samad

November 6<sup>th</sup>, 2024

# CONCURRENCY CONTROL SCHEMES

- One of the fundamental properties of a transaction is ***isolation***
- When several transactions execute concurrently in the database, the isolation property may no longer be preserved
- To ensure isolation, the system must control the interaction among the concurrent transactions
- This control is achieved through ***concurrency-control schemes***
- Cascadeless schedules overcome the temporary update problem, and ensures isolation property, as a new transaction doesn't start its operations until the previous transaction has completed its task

# CONCURRENCY CONTROL SCHEMES

- Example:
  - For the given schedule, check if it's in recoverable state or not? If not, get its cascadeless solution
- Solution:
  - On board

$T_1$	$T_2$
read( $B$ )	
$B := B - 50$	
write( $B$ )	
	read( $B$ )
	read( $A$ )
	display( $B + A$ )
	commit
read( $A$ )	
$A := A + 50$	
write( $A$ )	
commit	

# CONCURRENCY CONTROL SCHEMES

- We can achieve a cascadeless schedule of the example in previous slide, but how do we make the system behave such that  $T_2$  stops its execution until  $T_1$  has finished execution?
  - We use *locks*

# LOCKS

- The most common method to implement concurrency control is using a ***lock***
- A lock is a mechanism which provides access privileges to a transaction on a shared resource
- While the lock is acquired, no other transaction can access that resource.
- In order to release the resource, i.e. to make it available to other transactions, a transaction has to ***unlock*** the resource.

# LOCK MODES

- We present two modes in which a data item may be locked.
  - **Shared.** If a transaction  $T_i$  has obtained a ***shared-mode*** lock (***S***) on item  $Q$ , then  $T_i$  can read, but cannot write  $Q$ .
  - **Exclusive.** If a transaction  $T_i$  has obtained an ***exclusive-mode*** lock (***X***) on item  $Q$ , then  $T_i$  can both read and write  $Q$ .
- Every transaction must request to *the **concurrency-control manager***.
- The transaction can proceed with the operation only after the concurrency-control manager ***grants*** the lock to the transaction.

# LOCK COMPATIBILITY

- Given a set of lock modes, we can define a compatibility function on them as follows:
- Let  $A$  and  $B$  represent arbitrary lock modes.
- Suppose that a transaction  $T_i$  requests a lock of mode  $A$  on item  $Q$  on which transaction  $T_j$  ( $T_i \neq T_j$ ) currently holds a lock of mode  $B$ .
- If transaction  $T_i$  can be granted a lock on  $Q$  immediately, in spite of the presence of the mode  $B$  lock, then we say mode  $A$  is **compatible** with mode  $B$ .

# LOCK COMPATIBILITY

- Such a function can be represented conveniently by a matrix.
- The ***compatibility relation*** between the two modes of locking discussed in this section appears in the matrix comp below.
- An element ***comp***( $A, B$ ) of the matrix has the value true if and only if mode  $A$  is compatible with mode  $B$ .

	S	X
S	true	false
X	false	false



# TRANSACTIONS WITH LOCKS

- Consider the following transactions  $T_1$  and  $T_2$  from the example on previous slides
- The objective is to transfer 50 units from  $B$  to  $A$ , and then display  $B + A$ .
- To maintain consistency,  $B + A$  must be the same before and after the transactions. (Let  $A = 100$ ,  $B = 200$ , then  $A + B = 300$ ).

```
 $T_1$ : lock-X( $B$ );  
      read( $B$ );  
       $B := B - 50$ ;  
      write( $B$ );  
      unlock( $B$ );  
      lock-X( $A$ );  
      read( $A$ );  
       $A := A + 50$ ;  
      write( $A$ );  
      unlock( $A$ ).
```

```
 $T_2$ : lock-S( $B$ );  
      read( $B$ );  
      unlock( $B$ );  
      lock-S( $A$ );  
      read( $A$ );  
      unlock( $A$ );  
      display( $B + A$ ).
```

# TRANSACTIONS WITH LOCKS

- As done in the example in the earlier slides, one Interleaved/Concurrent schedule may look like the schedule given on the right
- We can introduce locks before every read/write operation

$T_1$	$T_2$
read( $B$ )	
$B := B - 50$	
write( $B$ )	
	read( $B$ )
	read( $A$ )
	display( $B + A$ )
	commit
read( $A$ )	
$A := A + 50$	
write( $A$ )	
commit	

# AN INCONSISTENT SCHEDULE

- The schedule with locks look like the schedule given on the right
- However, this does not resolve the temporary update problem
- Also, despite having the acquisition of locks the following schedule results in an inconsistent state
- This is because we simply added the locks without taking inconsistency problems into consideration, like lost update or temporary update or incorrect summary problems

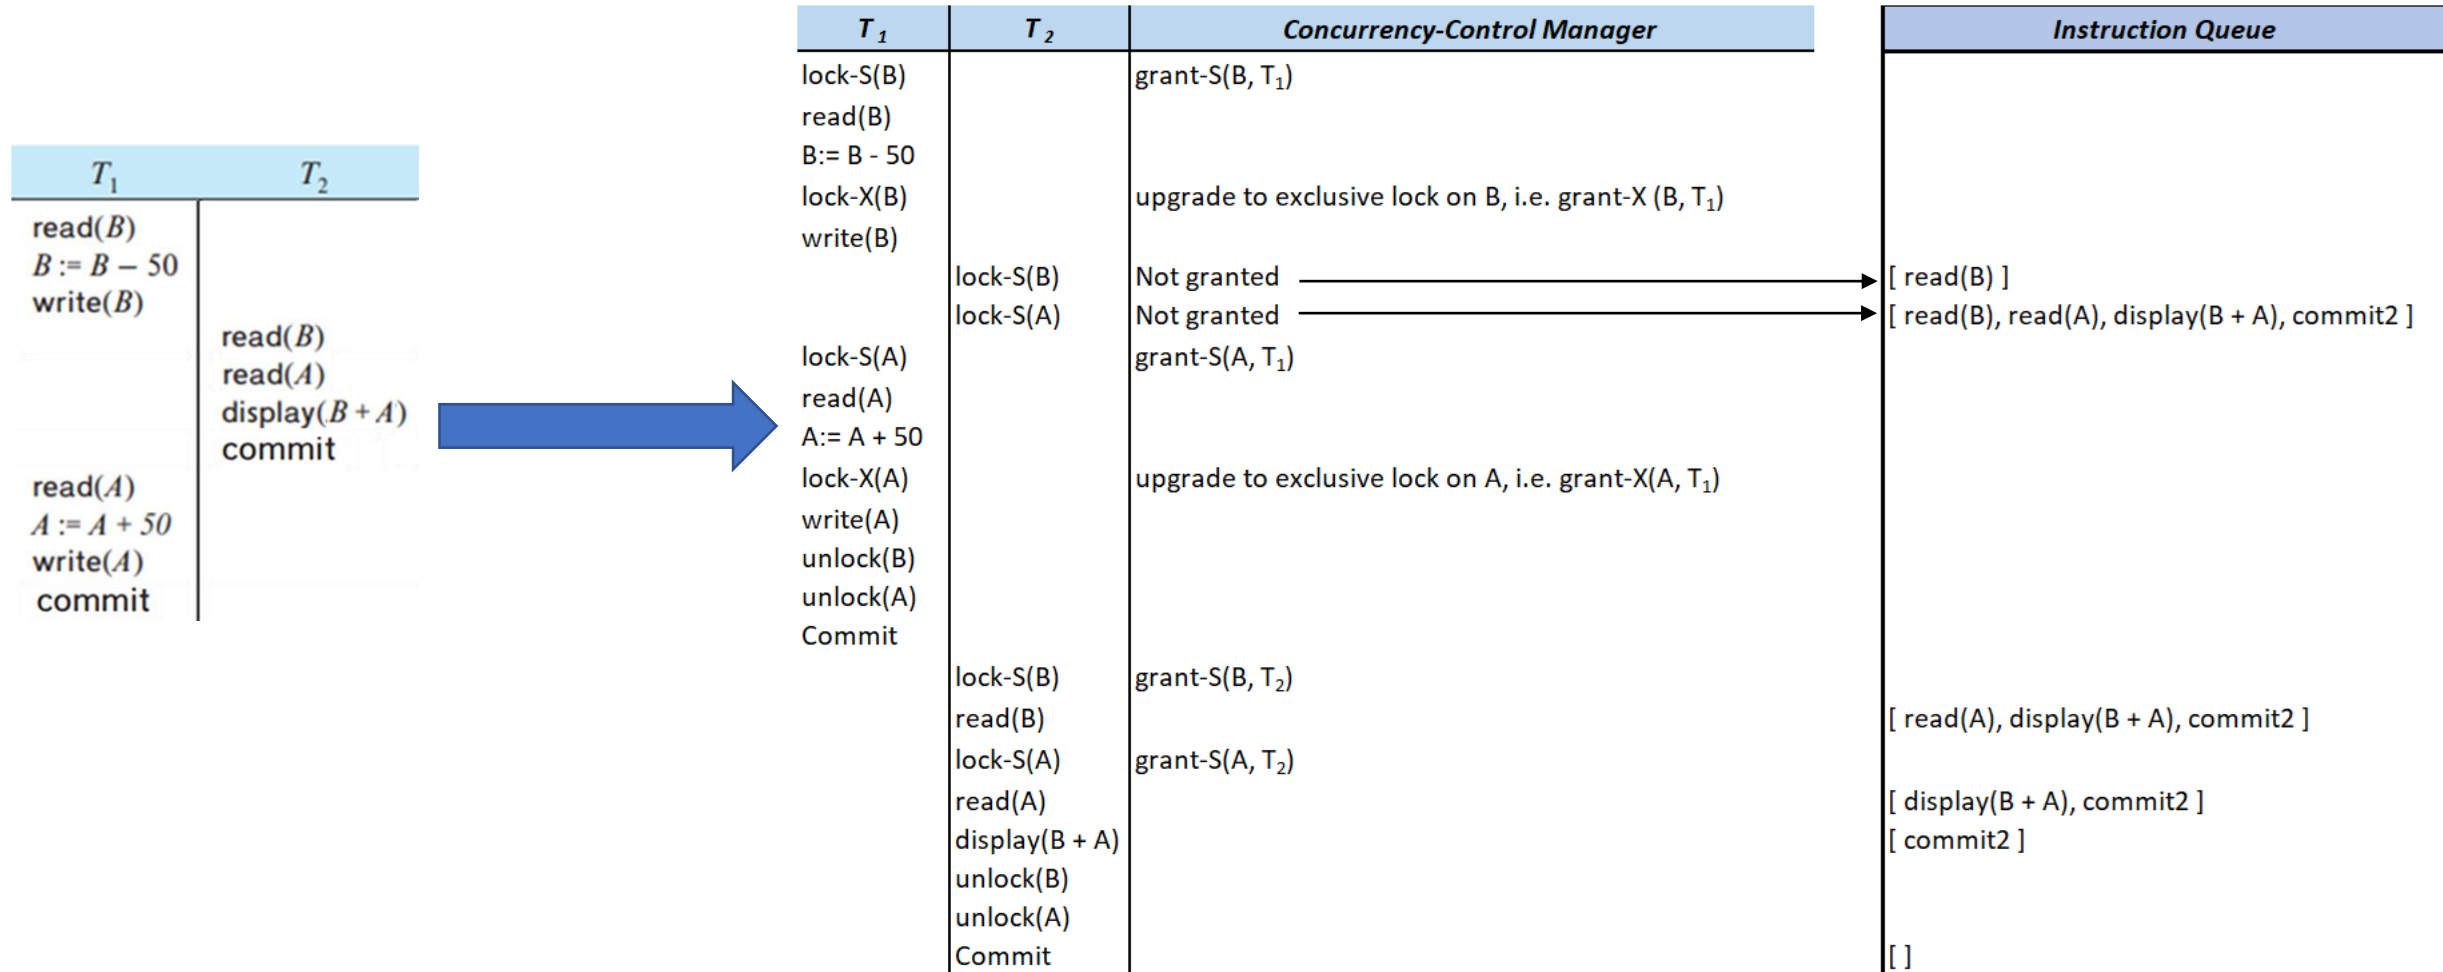
$T_1$	$T_2$	Concurrency-Control Manager
lock-S(B)		grant-S(B, $T_1$ )
read(B)		
$B := B - 50$		
lock-X(B)		upgrade to exclusive lock on B, i.e. grant-X(B, $T_1$ )
write(B)	B = 150	
unlock(B)		
	lock-S(B)	grant-S(B, $T_2$ )
	read(B)	
	unlock(B)	
	lock-S(A)	grant-S(A, $T_2$ )
	read(A)	
	unlock(A)	
	display(B + A)	
	Commit	
		grant-S(A, $T_1$ )
		upgrade to exclusive lock on A, i.e. grant-X(A, $T_1$ )
lock-S(A)		
read(A)		
$A := A + 50$		
lock-X(A)		
write(A)		
unlock(A)		
Commit		

A + B = 250

# A CONSISTENT SCHEDULE

- Update to have consistent schedule
- A simple rule is to **not** grant any type of lock to any transaction while the previous transaction is executing. A transaction “delays” the unlocking until it’s in a partially committed state
- Once the first transaction has written and released the lock, the second transaction may get access.
- For both transactions reading, that will not be an issue, as long as there is no write-read dependencies
- This forms a series of operations that resemble the solution of cascadeless schedules

# A CONSISTENT SCHEDULE



# DELAYED UNLOCKING – CONSISTENT SCHEDULE

- With unlocking delayed, transactions  $T_3$  and  $T_4$  lead to a consistent schedule.
- It also becomes a serial schedule, giving consistent results
- However, this may cause performance issues. For example, if  $T_3$  is being delayed due to IO operation,  $T_4$  would be idle for long time too

```
 $T_3$ : lock-X( $B$ );  
      read( $B$ );  
       $B := B - 50$ ;  
      write( $B$ );  
      lock-X( $A$ );  
      read( $A$ );  
       $A := A + 50$ ;  
      write( $A$ );  
      unlock( $B$ );  
      unlock( $A$ ).
```

```
 $T_4$ : lock-S( $A$ );  
      read( $A$ );  
      lock-S( $B$ );  
      read( $B$ );  
      display( $A + B$ );  
      unlock( $A$ );  
      unlock( $B$ ).
```

# DEADLOCKS

- Unfortunately, delayed locking can also lead to an undesirable situation.
- Consider the partial schedule as shown here for  $T_3$  and  $T_4$ .
- Since  $T_3$  is holding an exclusive-mode lock on  $B$  and  $T_4$  is requesting a shared-mode lock on  $B$ ,  $T_4$  is waiting for  $T_3$  to unlock  $B$ .
- Similarly, since  $T_4$  is holding a shared-mode lock on  $A$  and  $T_3$  is requesting an exclusive mode lock on  $A$ ,  $T_3$  is waiting for  $T_4$  to unlock  $A$ .
- Thus, we have arrived at a state where neither of these transactions can ever proceed with its normal execution. This situation is called **deadlock**.

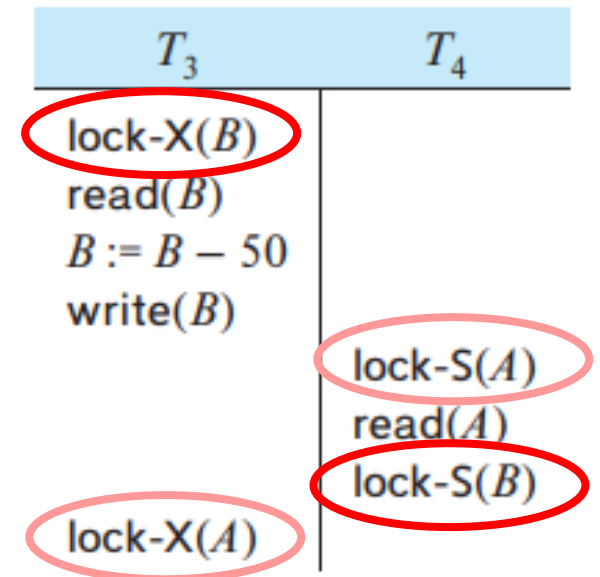



Figure 18.7 Schedule 2.

# DEADLOCK (TRAFFIC EXAMPLE)





# DEADLOCKS (EXAMPLE)

- Example:
  - S: r1(A), w1(A), r2(B), r3(B), r3(A), r1(B), w1(B), C1, r2(A), r3(A), C2, C3
  - Using the delayed locking strategy, check if the given schedule ends in a deadlock state or not?
- Solution:
  - On board

<u>T1</u>	<u>T2</u>	<u>T3</u>
r(A)		
w(A)		
	r(B)	
		r(B)
		r(A)
r(B)		
w(B)		
Commit		
	r(A)	
		r(A)
	Commit	
		Commit

# DEADLOCKS

- Activity Sheet

# DEADLOCKS

- Activity Sheet Solution:
  - Q1) Deadlock present
  - Q2) No deadlock and resolves fine