



# DATABASE SYSTEMS

CS – 355/CE – 373

Instructor: Maria N. Samad

October 23<sup>rd</sup>, 2024

# REVIEW – TRANSACTIONS

- A transaction (with its steps) must appear to the user as a single, indivisible unit. A transaction executes in its entirety or not at all.
- Several problems can occur due to concurrent execution, leading to the concepts of transaction processing.
- A transaction must have ACID properties: Atomicity, Consistency, Isolation, Durability.
- A transaction must be in one of the following states: Active, Partially Committed, Failed, Aborted or Committed.

# TRANSACTION SCHEDULES

- When transactions are executing concurrently in an interleaved fashion, then the order of execution of operations from all the various transactions is known as a **schedule** (or **history**).
- More formally, a **schedule** (or **history**)  $S$  of  $n$  transactions  $T_1, T_2, \dots, T_n$  is an ordering of the operations of the transactions.
- Operations from different transactions can be interleaved in the schedule  $S$ .
- However, for each transaction  $T_i$  that participates in the schedule  $S$ , the operations of  $T_i$  in  $S$  must appear in the same order in which they occur in  $T_i$ .

# SERIAL SCHEDULE

- A *serial schedule* consists of a sequence of instructions from various transactions, where the instructions belonging to one single transaction appear together in that schedule.

$T_1$	$T_2$
<code>read(A)</code> <code>A := A - 50</code> <code>write(A)</code> <code>read(B)</code> <code>B := B + 50</code> <code>write(B)</code> <code>commit</code>	<code>read(A)</code> <code>temp := A * 0.1</code> <code>A := A - temp</code> <code>write(A)</code> <code>read(B)</code> <code>B := B + temp</code> <code>write(B)</code> <code>commit</code>

$T_1$	$T_2$
<code>read(A)</code> <code>A := A - 50</code> <code>write(A)</code> <code>read(B)</code> <code>B := B + 50</code> <code>write(B)</code> <code>commit</code>	<code>read(A)</code> <code>temp := A * 0.1</code> <code>A := A - temp</code> <code>write(A)</code> <code>read(B)</code> <code>B := B + temp</code> <code>write(B)</code> <code>commit</code>

# CONCURRENT SCHEDULE – CONSISTENT STATE

$T_1$	$T_2$
read( $A$ ) $A := A - 50$ write( $A$ )	
	read( $A$ ) $temp := A * 0.1$ $A := A - temp$ write( $A$ )
read( $B$ ) $B := B + 50$ write( $B$ ) commit	
	read( $B$ ) $B := B + temp$ write( $B$ ) commit

**Figure 17.4** Schedule 3—a concurrent schedule equivalent to schedule 1.

# CONCURRENT SCHEDULE – INCONSISTENT STATE

$T_1$	$T_2$
$\text{read}(A)$ $A := A - 50$	$\text{read}(A)$ $\text{temp} := A * 0.1$ $A := A - \text{temp}$ $\text{write}(A)$ $\text{read}(B)$
$\text{write}(A)$ $\text{read}(B)$ $B := B + 50$ $\text{write}(B)$ $\text{commit}$	$B := B + \text{temp}$ $\text{write}(B)$ $\text{commit}$

**Figure 17.5** Schedule 4—a concurrent schedule resulting in an inconsistent state.

# SERIALIZABLE SCHEDULE

- A serial schedule (as shown in one of the previous slides) is always consistent, i.e. it maintains the consistent state of the database.
- However, the same cannot be guaranteed for a concurrent schedule.
- If a concurrent schedule can be shown to have the same effect as a serial schedule, (in other words it is shown to be equivalent to a serial schedule), then it can ensure the consistency of the database.
- Such a schedule is called a ***serializable schedule***.

# SERIALIZABLE SCHEDULES – EQUIVALENCE

$T_1$	$T_2$
$\text{read}(A)$ $A := A - 50$ $\text{write}(A)$	$\text{read}(A)$ $\text{temp} := A * 0.1$ $A := A - \text{temp}$ $\text{write}(A)$
$\text{read}(B)$ $B := B + 50$ $\text{write}(B)$ $\text{commit}$	$\text{read}(B)$ $B := B + \text{temp}$ $\text{write}(B)$ $\text{commit}$

**Figure 17.4** Schedule 3—a concurrent schedule equivalent to schedule 1.

E  
Q  
U  
I  
V  
A  
L  
E  
N  
T  
  
T  
O

$T_1$	$T_2$
$\text{read}(A)$ $A := A - 50$ $\text{write}(A)$ $\text{read}(B)$ $B := B + 50$ $\text{write}(B)$ $\text{commit}$	$\text{read}(A)$ $\text{temp} := A * 0.1$ $A := A - \text{temp}$ $\text{write}(A)$ $\text{read}(B)$ $B := B + \text{temp}$ $\text{write}(B)$ $\text{commit}$

**Figure 17.2** Schedule 1—a serial schedule in which  $T_1$  is followed by  $T_2$ .



# SERIALIZABLE SCHEDULE

- In the previous slide, it may seem we have simply combined the operations of each transactions and executed them one after the other. However, that may not be possible each time
- This happened because there was no dependency of data from one transaction to another.
- When there are such dependencies, we can't simply merge the operations and execute them in serial manner, as this may lead to inconsistent and/or incorrect output

# SERIALIZABLE SCHEDULES – EQUIVALENCE

$T_1$	$T_2$
read( $A$ ) $A := A - 50$	
	read( $A$ ) $A := A + 50$
write( $A$ ) commit	
	write( $A$ ) commit

I  
S  
  
N  
O  
T  
  
E  
Q  
U  
I  
V  
A  
L  
E  
N  
T  
  
T  
O

$T_1$	$T_2$
read( $A$ ) $A := A - 50$ write( $A$ ) commit	
	read( $A$ ) $A := A + 50$ write( $A$ ) commit

# SERIALIZABLE SCHEDULE

- So how do we know if a schedule is serializable or not?
- We achieve that by resolving conflicts between operations. If there are no conflicts, we can swap the operations with each other to get a resultant serial schedule
- We consider only *read(X)* and *write(X)* operations. The underlying assumption is that these two operations are the most significant operations from a scheduling perspective.
- All the other operations that happen between these significant operations will automatically be handled along with the reads and writes

# SERIALIZABLE SCHEDULES – EQUIVALENCE

- Lets consider a schedule  $S$  in which there are two consecutive instructions  $I$  and  $J$  of transactions  $T_i$  and  $T_j$  respectively.
- If  $I$  and  $J$  refer to different data items, then we can swap  $I$  and  $J$  without affecting the results of instructions in the schedule.
- However, if  $I$  and  $J$  refer to the same data item  $Q$ , then the order of the two steps may matter.

T1	T2
<i>read(X)</i>	
	<i>read(Y)</i>

Initial (Above), Swapped (Below)

T1	T2
	<i>read(Y)</i>
<i>read(X)</i>	

# SERIALIZABLE SCHEDULES – EQUIVALENCE

- Since we are dealing with only *read* and *write* instructions, there are four cases that we need to consider:
- **[CASE 1]  $I = \text{read}(Q), J = \text{read}(Q)$ :**
  - The order of  $I$  and  $J$  does not matter, since the same value of  $Q$  is read by  $T_i$  and  $T_j$ , regardless of the order.

T1	T2
<i>read(Q)</i>	
	<i>read(Q)</i>

Initial (Left),  
Swapped (Right)  
Equivalent

T1	T2
	<i>read(Q)</i>
<i>read(Q)</i>	

# SERIALIZABLE SCHEDULES – EQUIVALENCE

- **[CASE 2]  $I = \text{read}(Q), J = \text{write}(Q)$ :**

- If  $I$  comes before  $J$ , then  $T_i$  does not read the value of  $Q$  that is written by  $T_j$  in instruction  $J$ .
- If  $J$  comes before  $I$ , then  $T_i$  reads the value of  $Q$  that is written by  $T_j$ . Thus, the order of  $I$  and  $J$  matters.

T1	T2
<i>read(Q)</i>	
	<i>write(Q)</i>

Initial (Left),  
Swapped (Right)  
Not Equivalent

T1	T2
	<i>write(Q)</i>
<i>read(Q)</i>	

# SERIALIZABLE SCHEDULES – EQUIVALENCE

- [CASE 3]  $I = \text{write}(Q)$ ,  $J = \text{read}(Q)$ :
  - The order of  $I$  and  $J$  matters for reasons similar to those of the previous case.

T1	T2
<i>write(Q)</i>	
	<i>read(Q)</i>

Initial (Left),  
Swapped (Right)  
Not Equivalent

T1	T2
	<i>read(Q)</i>
<i>write(Q)</i>	

# SERIALIZABLE SCHEDULES – EQUIVALENCE

- [CASE 4]  $I = \text{write}(Q), J = \text{write}(Q)$ :
  - Since both instructions are write operations, the order of these instructions does not affect either  $T_i$  or  $T_j$ .
  - However, the value obtained by the next  $\text{read}(Q)$  instruction of  $S$  is affected, since the result of only the latter of the two write instructions is preserved in the database.
  - If there is no other  $\text{write}(Q)$  instruction after  $I$  and  $J$  in  $S$ , then the order of  $I$  and  $J$  directly affects the final value of  $Q$  in the database state that results from schedule  $S$ .

T1	T2
$\text{write}(Q)$	
	$\text{write}(Q)$

Initial (Left),  
Swapped (Right)  
Equivalent  
ONLY when there is no  
 $\text{read}(Q)$  in between

T1	T2
	$\text{write}(Q)$
$\text{write}(Q)$	



# SERIALIZABLE SCHEDULES – CONFLICTS

- Thus, only in the case where both  $I$  and  $J$  are read instructions does the relative order of their execution not matter.
- We say that  $I$  and  $J$  conflict if they are operations by different transactions on the same data item, and **at least one of these instructions is a write operation.**

Initial Schedule	Swapped Schedule	Conflict?												
<table><tr><th>T1</th><th>T2</th></tr><tr><td><i>read(Q)</i></td><td></td></tr><tr><td></td><td><i>read(Q)</i></td></tr></table>	T1	T2	<i>read(Q)</i>			<i>read(Q)</i>	<table><tr><th>T1</th><th>T2</th></tr><tr><td></td><td><i>read(Q)</i></td></tr><tr><td><i>read(Q)</i></td><td></td></tr></table>	T1	T2		<i>read(Q)</i>	<i>read(Q)</i>		No
T1	T2													
<i>read(Q)</i>														
	<i>read(Q)</i>													
T1	T2													
	<i>read(Q)</i>													
<i>read(Q)</i>														
<table><tr><th>T1</th><th>T2</th></tr><tr><td><i>read(Q)</i></td><td></td></tr><tr><td></td><td><i>write(Q)</i></td></tr></table>	T1	T2	<i>read(Q)</i>			<i>write(Q)</i>	<table><tr><th>T1</th><th>T2</th></tr><tr><td></td><td><i>write(Q)</i></td></tr><tr><td><i>read(Q)</i></td><td></td></tr></table>	T1	T2		<i>write(Q)</i>	<i>read(Q)</i>		Yes
T1	T2													
<i>read(Q)</i>														
	<i>write(Q)</i>													
T1	T2													
	<i>write(Q)</i>													
<i>read(Q)</i>														
<table><tr><th>T1</th><th>T2</th></tr><tr><td><i>write(Q)</i></td><td></td></tr><tr><td></td><td><i>read(Q)</i></td></tr></table>	T1	T2	<i>write(Q)</i>			<i>read(Q)</i>	<table><tr><th>T1</th><th>T2</th></tr><tr><td></td><td><i>read(Q)</i></td></tr><tr><td><i>write(Q)</i></td><td></td></tr></table>	T1	T2		<i>read(Q)</i>	<i>write(Q)</i>		Yes
T1	T2													
<i>write(Q)</i>														
	<i>read(Q)</i>													
T1	T2													
	<i>read(Q)</i>													
<i>write(Q)</i>														
<table><tr><th>T1</th><th>T2</th></tr><tr><td><i>write(Q)</i></td><td></td></tr><tr><td></td><td><i>write(Q)</i></td></tr></table>	T1	T2	<i>write(Q)</i>			<i>write(Q)</i>	<table><tr><th>T1</th><th>T2</th></tr><tr><td></td><td><i>write(Q)</i></td></tr><tr><td><i>write(Q)</i></td><td></td></tr></table>	T1	T2		<i>write(Q)</i>	<i>write(Q)</i>		Yes
T1	T2													
<i>write(Q)</i>														
	<i>write(Q)</i>													
T1	T2													
	<i>write(Q)</i>													
<i>write(Q)</i>														

# SERIALIZABLE SCHEDULES – SWAPPING

- Let  $I$  and  $J$  be consecutive instructions of a schedule  $S$ .
- If  $I$  and  $J$  are instructions of different transactions and  $I$  and  $J$  do not conflict, then we can swap the order of  $I$  and  $J$  to produce a new schedule  $S'$ .
- $S$  is equivalent to  $S'$ , since all instructions appear in the same order in both schedules except for  $I$  and  $J$ , whose order does not matter.

# SWAPPING

- Consider the following schedule →
- We can swap the instructions *read(B)* and *write(A)* as they do not conflict.

$T_1$	$T_2$
read( $A$ ) write( $A$ )	
	read( $A$ ) write( $A$ )
read( $B$ ) write( $B$ )	
	read( $B$ ) write( $B$ )

# SWAPPING

- The resulting schedule is shown →
- We can again swap the instructions *read(B)* and *read(A)* as they do not conflict.

$T_1$	$T_2$
read( $A$ )	
write( $A$ )	
read( $B$ )	read( $A$ )
	write( $A$ )
write( $B$ )	read( $B$ )
	write( $B$ )

# SWAPPING

- The resulting schedule is shown →
- We can now swap the instructions *write(B)* and *write(A)* as they do not conflict.

$T_1$	$T_2$
read( $A$ ) write( $A$ ) read( $B$ )	
	read( $A$ ) write( $A$ )
write( $B$ )	
	read( $B$ ) write( $B$ )

# SWAPPING

- The resulting schedule is shown →
- Finally, we can swap the instructions  $write(B)$  and  $read(A)$  as they do not conflict.

$T_1$	$T_2$
read( $A$ ) write( $A$ ) read( $B$ )	
write( $B$ )	read( $A$ )
	write( $A$ ) read( $B$ ) write( $B$ )

# SWAPPING

- The resulting schedule is shown →
- Note that this schedule is a serial schedule.
- Since the original schedule has been shown to be equivalent to a serial schedule, we conclude that the original schedule maintains the consistency of the database.

$T_1$	$T_2$
read( $A$ ) write( $A$ ) read( $B$ ) write( $B$ )	read( $A$ ) write( $A$ ) read( $B$ ) write( $B$ )

# CONFLICT EQUIVALENT SCHEDULES

- If a schedule  $S$  can be transformed into a schedule  $S'$  by a series of swaps of nonconflicting instructions, we say that  $S$  and  $S'$  are **conflict equivalent**.

$T_1$	$T_2$	E Q U I V A L E N T  T O	$T_1$	$T_2$
read( $A$ )			read( $A$ )	
write( $A$ )			write( $A$ )	
	read( $A$ )		read( $B$ )	
	write( $A$ )		write( $B$ )	
read( $B$ )				read( $A$ )
write( $B$ )				write( $A$ )
	read( $B$ )			read( $B$ )
	write( $B$ )			write( $B$ )



# CONFLICT EQUIVALENT SCHEDULES

- Example: Given three schedules as follows:
  - S1: r1(X), w1(X), r2(X), w2(X), r1(Y), w1(Y), r2(Y), w2(Y)
  - S2: r1(X), w1(X), r1(Y), r2(X), w2(X), w1(Y), r2(Y), w2(Y)
- Check if the two schedules are conflict equivalent to each other by using swapping techniques
- Solution: On board

# CONFLICT EQUIVALENT SCHEDULES

- Example: Given two schedules as follows:
  - S1:  $r_1(X)$ ,  $r_2(Y)$ ,  $w_2(Y)$ ,  $w_1(X)$ ,  $w_3(X)$
  - S2:  $r_1(X)$ ,  $w_1(X)$ ,  $r_2(Y)$ ,  $w_3(X)$ ,  $w_2(Y)$
- Check if the schedules are conflict equivalent to each other by using swapping techniques
- Solution: On board

# CONFLICT EQUIVALENT SCHEDULES

- Example: Given two schedules as follows:
  - S1:  $r_1(X)$ ,  $r_2(X)$ ,  $w_1(X)$ ,  $w_2(X)$ ,  $w_3(X)$
  - S2:  $r_1(X)$ ,  $w_1(X)$ ,  $r_2(X)$ ,  $w_2(X)$ ,  $w_3(X)$
- Check if the schedules are conflict equivalent to each other by using swapping techniques
- Solution: On board

# CONFLICT EQUIVALENT SCHEDULES

- Class Activity

# CONFLICT EQUIVALENT SCHEDULES

- Class Activity Solution
  - [Conflict Equivalent Schedules Solution](#)