



DATABASE SYSTEMS

CS – 355/CE – 373

Instructor: Maria N. Samad

November 4th, 2024

RECOVERABLE SCHEDULE

- Even if the schedule is recoverable, there can be schedules that do not have single dependencies, or even have cyclic dependencies
- A schedule that has transitive dependencies is known to be ***cascaded***, i.e. series of dependencies that rely on the previous ones.
- Cascaded dependencies are not single dependencies, and must be handled in depth rather than only checking for the order of commit operations
 - Example: $T_1 \rightarrow T_2$ and $T_2 \rightarrow T_3$ and $T_3 \rightarrow T_4$
- A schedule has cyclic dependency when the transactions depend on each other
 - Example: $T_1 \rightarrow T_2$ and $T_2 \rightarrow T_1$

RECOVERABLE SCHEDULE

- For cascaded transactions, we deduce what is commonly known as the ***cascadeless schedule***.
- This method goes beyond just moving the commit operations to deduce it is “recoverable” or not
- Instead it converts the schedule in a form that prevents inconsistencies occurring due to failure in between, i.e. the temporary update problem

RECOVERABLE SCHEDULE

- However, schedules containing cyclic write – read dependencies are ALWAYS ***non-recoverable***, and cannot be converted into recoverable schedules either
- Thus, for cyclic dependencies, there is no need to move the commit operations or deduce cascadeless schedule, and we can directly say that if a failure occurs in such schedules, they can NEVER be converted into a recoverable state.
- The only possible option can be rewriting the schedule as a compensating transaction by removing the cycle

CASCADING ROLLBACK

- Even if a schedule is recoverable, to recover correctly from the failure of a transaction T_i , we may have to roll back several transactions.
- As an illustration, consider the partial schedule shown here
→
- Transaction T_8 writes a value of A that is read by transaction T_9 . Transaction T_9 writes a value of A that is read by transaction T_{10} .

| T_8 | T_9 | T_{10} |
|--|-----------------------------|-------------|
| read(A) read(B) write(A) | read(A) write(A) | read(A) |
| abort | | |

CASCADING ROLLBACK

- Suppose that, at this point, T_8 fails. T_8 must be rolled back.
- Since T_9 is dependent on T_8 , T_9 must be rolled back.
- Since T_{10} is dependent on T_9 , T_{10} must be rolled back.
- This phenomenon, in which a single transaction failure leads to a series of transaction rollbacks, is called **cascading rollback**

| T_8 | T_9 | T_{10} |
|--|-----------------------------|-------------|
| read(A) read(B) write(A) | read(A) write(A) | read(A) |
| abort | | |

CASCADELESS SCHEDULES

- Cascading rollback is undesirable, since it leads to the undoing of a significant amount of work.
- It is desirable to restrict the schedules to those where cascading rollbacks cannot occur.
- Such schedules are called cascadeless schedules.
- Formally, a cascadeless schedule is one where, for each pair of transactions T_i and T_j such that T_j reads a data item previously written by T_i , the commit operation of T_i appears before the read operation of T_j .
- **Every cascadeless schedule is also recoverable.**

EXAMPLE

- Consider the following schedule and deduce the following:
 1. Are there any cascading rollbacks?
 2. Is the schedule in recoverable state or not?
 3. If not, what is the cascadeless schedule for it?
- Solution:
 - Write – Read dependencies:
 - $T1 \rightarrow T2 (A)$
 - $T2 \rightarrow T3 (A)$
 - Check $T1 \rightarrow T3$ for possibility of overall cycles?
 - $T1 \rightarrow T3$ exists for A
 - $T3 \rightarrow T1 (B)$
 - $T3 \rightarrow T2 (B)$
 - 1. Cascading rollback: $T1 \rightarrow T2 \rightarrow T3 (A)$
 - 2. No, because cycles present:
 - $T1 \rightarrow T3 (A)$ & $T3 \rightarrow T1 (A)$
 - $T2 \rightarrow T3 (B)$ & $T3 \rightarrow T2 (B)$
 - 3. Can't deduce cascadeless schedule because it cannot be recovered due to cyclic dependencies

| <u>T1</u> | <u>T2</u> | <u>T3</u> |
|-----------|-----------|-----------|
| w(A) | | |
| | r(A) | |
| r(B) | | |
| | | w(B) |
| | | r(A) |
| | w(A) | |
| r(B) | | |
| | | r(A) |
| | | w(A) |
| | r(B) | |
| Commit | | |
| | | r(B) |
| | w(B) | |
| | | Commit |
| | Commit | |

EXAMPLE

- Consider the following schedule and deduce the following:
 1. Are there any cascading rollbacks?
 2. Is the schedule in recoverable state or not?
 3. If not, what is the cascadeless schedule for it?
- Solution:
 - Remember, a cascadeless schedule is by default recoverable schedule
 - Write – Read dependencies:
 - $T1 \rightarrow T2 (A)$
 - $T2 \rightarrow T3 (A)$
 - Check $T1 \rightarrow T3$ for possibility of overall cycles?
 - $T1 \rightarrow T3$ exists for A
 - $T2 \rightarrow T3 (A)$
 - Check $T1 \rightarrow T3$ for possibility of overall cycles?
 - $T1 \rightarrow T3$ exists for A
 - 1. Cascading rollback: $T1 \rightarrow T2 \rightarrow T3$
 - 2. No cycles, so it might possibly be recoverable
 - On board
 - 3. Recoverable, so get cascadeless schedule
 - On board

| <u>T1</u> | <u>T2</u> | <u>T3</u> |
|-----------|-----------|-----------|
| w(A) | | |
| | r(A) | |
| r(B) | | |
| | w(A) | |
| | | r(A) |
| | w(A) | |
| r(B) | | |
| | | r(A) |
| | | w(A) |
| | r(B) | |
| Commit | | |
| | | r(B) |
| | w(B) | |
| | | Commit |
| | Commit | |

CASCADELESS SCHEDULE

- Activity Sheet

CASCADELESS SCHEDULE

- Activity Sheet Solution:
 - [Cascadeless Schedules Solution](#)